

```
In [16]: import numpy as np
import pandas as pd
from sklearn.datasets import make_classification
```

```
In [17]: X, y = make_classification(n_samples=50000, n_features=15, n_informativ
e=10, n_redundant=5,
                                n_classes=2, weights=[0.7], class_sep=0.7, r
andom_state=15)
```

```
In [18]: X.shape, y.shape
```

```
Out[18]: ((50000, 15), (50000,))
```

```
In [19]: from sklearn.model_selection import train_test_split
```

```
In [20]: X_train, X_test, y_train, y_test = train_test_split(X, y, test_size=0.2
5, random_state=15)
```

```
In [21]: X_train.shape, y_train.shape, X_test.shape, y_test.shape
```

```
Out[21]: ((37500, 15), (37500,)), (12500, 15), (12500,))
```

```
In [22]: from sklearn import linear_model
```

```
In [23]: # alpha : float
# Constant that multiplies the regularization term.

# eta0 : double
# The initial learning rate for the 'constant', 'invscaling' or 'adaptive'
schedules.
```

```
clf = linear_model.SGDClassifier(eta0=0.0001, alpha=0.0001, loss='log',
    random_state=15, penalty='l2', tol=1e-3, verbose=2, learning_rate='constant')
clf
```

```
Out[23]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
    early_stopping=False, epsilon=0.1, eta0=0.0001,
    fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
    loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
    penalty='l2', power_t=0.5, random_state=15, shuffle=True,
    tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [24]: clf.fit(X=X_train, y=y_train)
```

```
-- Epoch 1
Norm: 0.77, NNZs: 15, Bias: -0.316653, T: 37500, Avg. loss: 0.455552
Total training time: 0.03 seconds.
-- Epoch 2
Norm: 0.91, NNZs: 15, Bias: -0.472747, T: 75000, Avg. loss: 0.394686
Total training time: 0.05 seconds.
-- Epoch 3
Norm: 0.98, NNZs: 15, Bias: -0.580082, T: 112500, Avg. loss: 0.385711
Total training time: 0.08 seconds.
-- Epoch 4
Norm: 1.02, NNZs: 15, Bias: -0.658292, T: 150000, Avg. loss: 0.382083
Total training time: 0.09 seconds.
-- Epoch 5
Norm: 1.04, NNZs: 15, Bias: -0.719528, T: 187500, Avg. loss: 0.380486
Total training time: 0.11 seconds.
-- Epoch 6
Norm: 1.05, NNZs: 15, Bias: -0.763409, T: 225000, Avg. loss: 0.379578
Total training time: 0.12 seconds.
-- Epoch 7
Norm: 1.06, NNZs: 15, Bias: -0.795106, T: 262500, Avg. loss: 0.379150
Total training time: 0.13 seconds.
-- Epoch 8
```

```
Norm: 1.06, NNZs: 15, Bias: -0.819925, T: 300000, Avg. loss: 0.378856
Total training time: 0.15 seconds.
-- Epoch 9
Norm: 1.07, NNZs: 15, Bias: -0.837805, T: 337500, Avg. loss: 0.378585
Total training time: 0.16 seconds.
-- Epoch 10
Norm: 1.08, NNZs: 15, Bias: -0.853138, T: 375000, Avg. loss: 0.378630
Total training time: 0.18 seconds.
Convergence after 10 epochs took 0.18 seconds
```

```
Out[24]: SGDClassifier(alpha=0.0001, average=False, class_weight=None,
                      early_stopping=False, epsilon=0.1, eta0=0.0001,
                      fit_intercept=True, l1_ratio=0.15, learning_rate='constant',
                      loss='log', max_iter=1000, n_iter_no_change=5, n_jobs=None,
                      penalty='l2', power_t=0.5, random_state=15, shuffle=True,
                      tol=0.001, validation_fraction=0.1, verbose=2, warm_start=False)
```

```
In [25]: clf.coef_, clf.coef_.shape, clf.intercept_
```

```
Out[25]: (array([[ -0.42336692,  0.18547565, -0.14859036,  0.34144407, -0.2081867
,
                0.56016579, -0.45242483, -0.09408813,  0.2092732 ,  0.1808412
6,
                0.19705191,  0.00421916, -0.0796037 ,  0.33852802,  0.0226672
1]]),
          (1, 15),
          array([ -0.8531383]))
```

## Implement Logistic Regression with L2 regularization Using SGD: without using sklearn

## Instructions

- Load the datasets(train and test) into the respective arrays
- Initialize the weight\_vector and intercept term randomly
- Calculate the initial log loss for the train and test data with the current weight and intercept and store it in a list
- for each epoch:
  - for each batch of data points in train: (keep batch size=1)
    - calculate the gradient of loss function w.r.t each weight in weight vector
    - Calculate the gradient of the intercept [check this](#)
    - Update weights and intercept (check the equation number 32 in the above mentioned [pdf](#)):
$$w^{(t+1)} \leftarrow (1 - \frac{\alpha \lambda}{N})w^{(t)} + \alpha x_n (y_n - \sigma((w^{(t)})^T x_n + b^t))$$
$$b^{(t+1)} \leftarrow (b^t + \alpha (y_n - \sigma((w^{(t)})^T x_n + b^t)))$$
    - calculate the log loss for train and test with the updated weights (you can check the python assignment 10th question)
    - And if you wish, you can compare the previous loss and the current loss, if it is not updating, then you can stop the training
    - append this loss in the list ( this will be used to see how loss is changing for each epoch after the training is over )
- Plot the train and test loss i.e on x-axis the epoch number, and on y-axis the loss
- **GOAL:** compare your implementation and SGDClassifier's the weights and intercept, make sure they are as close as possible i.e difference should be in terms of  $10^{-3}$

```
In [26]: X, y = make_classification(n_samples=50000, n_features=15, n_informativ
```

```
e=10, n_redundant=5,  
                                n_classes=2, weights=[0.7], class_sep=0.7, r  
andom_state=15)
```

```
In [27]: w = np.zeros_like(X_train[0]) #weight vector  
b = 0 #intercept term  
eta0 = 0.0001  
alpha = 0.0001  
N = len(X_train)  
#print(w)
```

```
In [28]: def sigmoid(w,x,b):  
        return 1/(1+np.exp(-(np.dot(x,w.T)+b)))
```

```
In [29]: #referenced from https://www.kaggle.com/anandkenta/logistic-regression-  
l2-regularization-using-sgd  
from tqdm import tqdm  
def log_function(w,x,y,b):  
    log_func_val = sigmoid(w,x,b)  
    a = y*np.log10(log_func_val)  
    b = (1-y)*np.log10(1-log_func_val)  
    output = -a-b  
    return np.mean(output)  
def trainSGD(x_train,y_train,x_test,y_test,eta0,alpha,epochs,w,b):  
    log_func_t = 1  
    log_loss_val_test = []  
    log_loss_val_train = []  
    for epoch in range(0,epochs):  
        for i in range(0,len(x_train)):  
            y = y_train[i]  
            x = x_train[i]  
            w = ((1-eta0*(alpha/N))*w)+((eta0*x)*(y-sigmoid(w,x,b)))  
            b = b+(eta0*(y-sigmoid(w,x,b)))  
  
            log_func_train = log_function(w, x_train, y_train,b)  
            log_loss_val_train.append(log_func_train)  
  
            log_func_test = log_function(w, x_test, y_test,b)
```

```

log_loss_val_test.append(log_func_test)

return w, b, log_loss_val_train, log_loss_val_test

```

```

In [30]: epochs = 10
        w1,b1,logLoss,logLoss_test = trainSGD(X_train,y_train,X_test,y_test, e
        ta0,alpha, epochs, w, b)
        print(w1,b1)
        print(logLoss)
        print(logLoss_test)

```

```

[-0.42315311  0.19095979 -0.14588118  0.33814991 -0.21196623  0.5652597
 8
 -0.44538357 -0.09171679  0.21795314  0.16977398  0.19522044  0.0022955
 4
 -0.07781461  0.33882618  0.02214234] -0.8500967712837225
[0.17546926223702466, 0.16868174436540248, 0.16639953379688374, 0.16537
404901928135, 0.16486122004082468, 0.16459114506307726, 0.1644447987447
564, 0.1643641152252568, 0.16431912310828212, 0.16429382915597823]
[0.17596687861916305, 0.1694098961177956, 0.1672141530442447, 0.1662232
9469756666, 0.16572403546384049, 0.16545876819806654, 0.165313652220769
9, 0.16523283564551378, 0.16518727864511057, 0.1651613611606301]

```

```

In [36]: w1-clf.coef_#b1-clf.intercept_ #difference between custom updated weigh
        ts

```

```

Out[36]: array([[ 0.0002138 ,  0.00548413,  0.00270918, -0.00329416, -0.0037795
 3,
                0.00509399,  0.00704126,  0.00237134,  0.00867994, -0.0110672
 8,
                -0.00183147, -0.00192361,  0.00178909,  0.00029817, -0.0005248
 7]])

```

```

In [37]: b1-clf.intercept_#difference between custom updated intercept term and
        clf.intercept_

```

```

Out[37]: array([0.00304153])

```

### observation:

Hence the result is matched with a sklearn

```
In [32]: def pred(w,b, X):
          N = len(X)
          predict = []
          #sigmoid = 0
          for i in range(N):
              if sigmoid(w, X[i], b) >= 0.5: # sigmoid(w,x,b) returns 1/(1+exp(-(dot(x,w)+b)))
                  predict.append(1)
              else:
                  predict.append(0)
          return np.array(predict)
          print(1-np.sum(y_train - pred(w,b,X_train))/len(X_train))
          print(1-np.sum(y_test - pred(w,b,X_test))/len(X_test))
```

```
1.6978933333333335
1.6986400000000001
```

```
In [33]: from sklearn.metrics import accuracy_score
          print(accuracy_score(y_train, pred(w1,b1,X_train)))
          print(accuracy_score(y_test, pred(w1,b1,X_test)))
```

```
0.83136
0.83376
```

```
In [35]: import matplotlib.pyplot as plt
          plt.plot(np.array(logLoss), color='darkblue', label='Log loss Training Data')
          plt.plot(np.array(logLoss_test), color='orange', label='Log loss Test Data')

          plt.ylabel('LogLoss')
          plt.xlabel('epochs')
          plt.title('Log loss curve')
```

```
plt.legend()  
plt.show()
```

